

Class introduction

APSC 8280: Machine learning applied to plant science

Outline

- Introductions
- Syllabus
- Introduction to Machine Learning
- Brief History of Machine Learning
- Lab 0: laptop setup
- Basic R programming I

Introductions

- Name
- Research
- Expectations from this class
- Something about you no one would have ever guessed

Syllabus: what to expect from this class

The basics

History of ML
Fundamentals of ML
Strengths and weaknesses of different ML models

Asking the right questions

Data cleaning and preprocessing
Feature engineering

Applied ML

Quantitative genetics
Plant disease epidemiology
Precision agriculture

Advanced topics

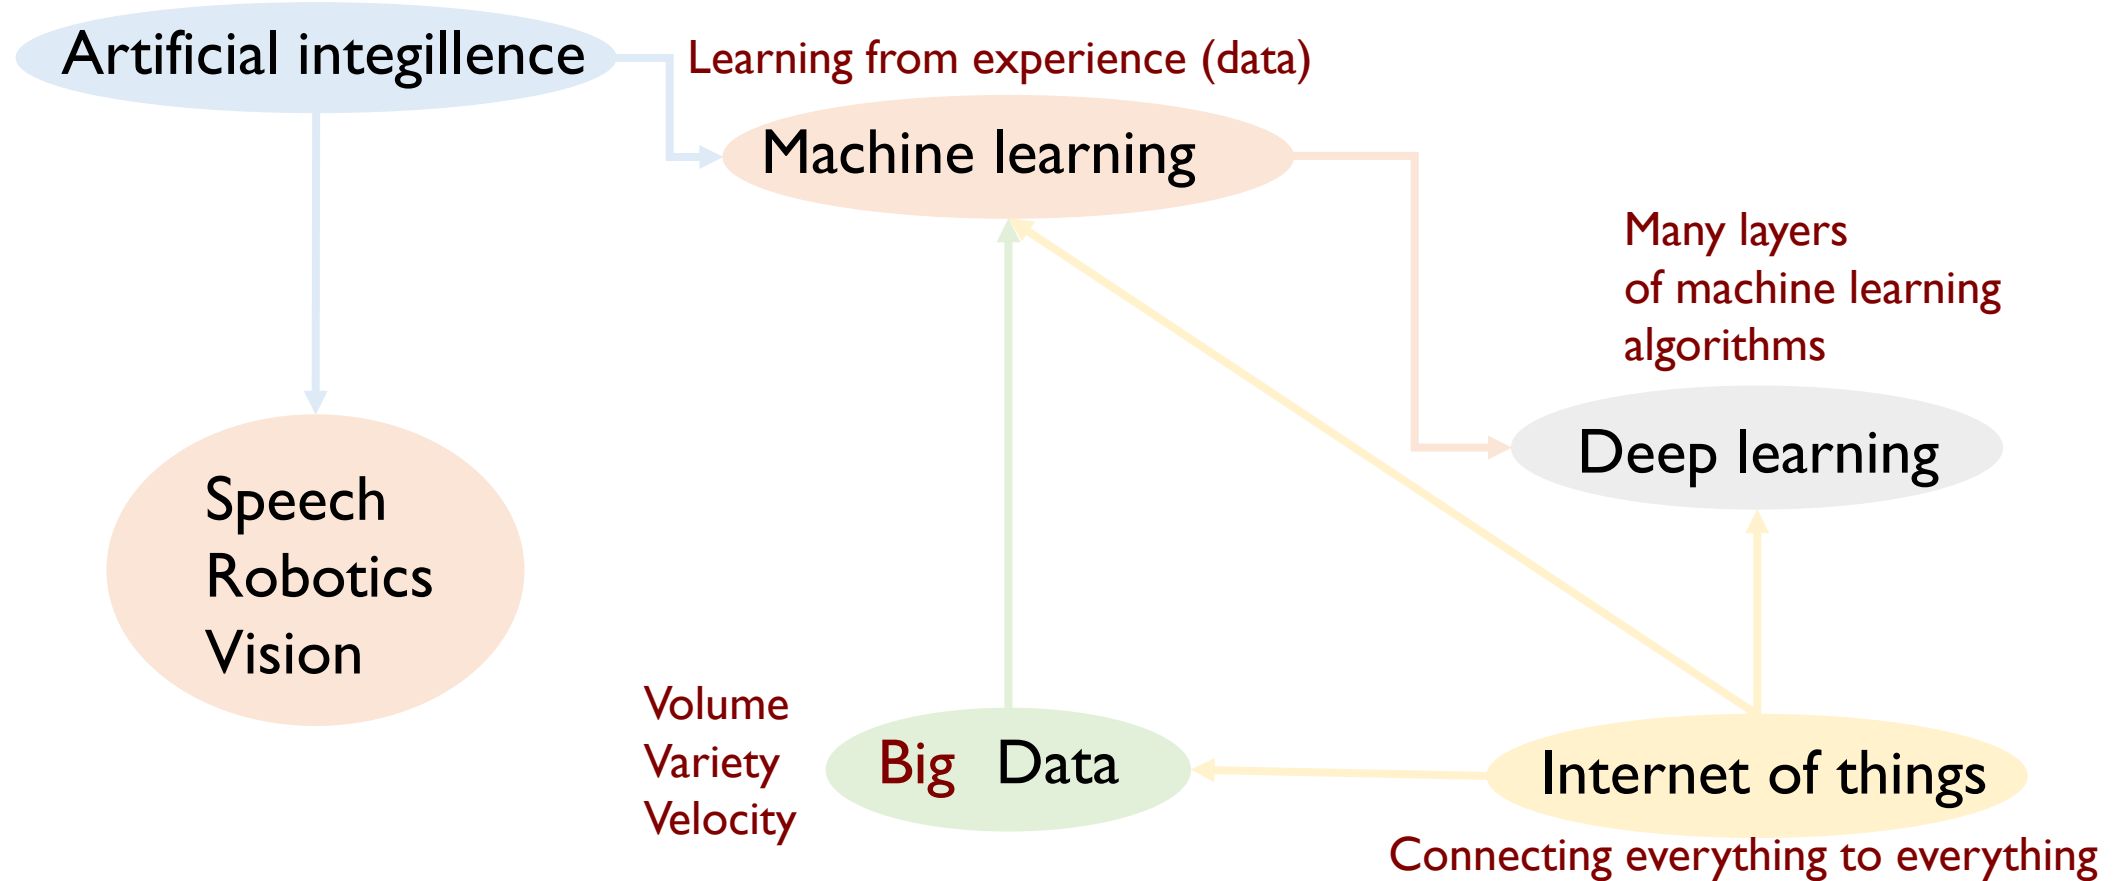
Deep learning
Computer vision
Reinforcement learning

What is not covered?

- Math behind machine learning algorithms and optimization
- Working with cloud and enterprise machine learning tools
- Deployment and maintenance of machine learning systems


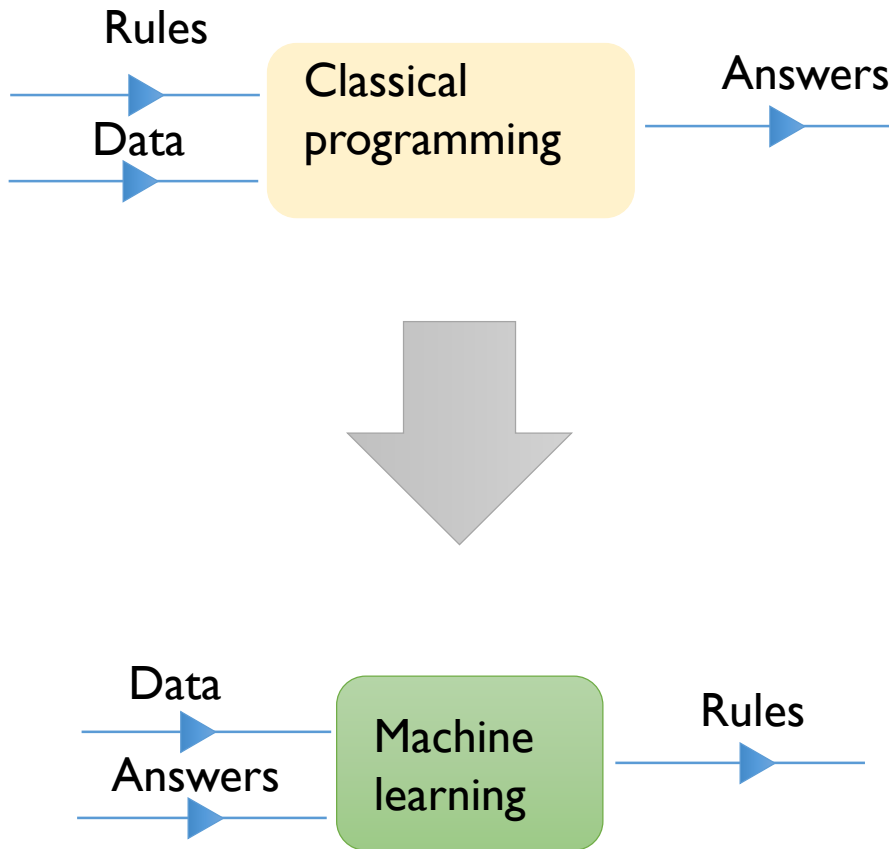
What's what?

Getting computers to do things humans do



Machine learning: a paradigm shift


“The unreasonable effectiveness of data”




- Banko, Michele, and Eric Brill. "Scaling to very very large corpora for natural language disambiguation." Proceedings of the 39th annual meeting on association for computational linguistics. Association for Computational Linguistics, 2001.
- Adapted from Géron, Aurélien. Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems. O'Reilly Media, 2019.

Types of machine learning


Supervised learning



Unsupervised learning



Reinforcement learning



Some tasks best solved by learning

Prediction

- Future stock prices or currency exchange rates
- Which movies will a person like
- The performance of crossing any two inbred and tester parents

Recognizing patterns

- Objects in real scenes
- Facial identities or facial expressions
- Spoken words

Recognizing anomalies

- Unusual sequences of credit card transactions
- Unusual patterns of sensor readings in a nuclear power plant

What has machine learning achieved so far?

Human-like skills

- ✓ Image classification
- ✓ Speech transcription
- ✓ Handwriting transcription
- ✓ Autonomous driving

Commerce

- ✓ Ad targeting
- ✓ Search results
- ✓ Customer service
- ✓ Digital assistants (e.g. Alexa)
- ✓ Reduction in fraudulent credit card transactions
- ✓ Stock market and investing

Healthcare

- ✓ Personalized/precision medicine
- ✓ Drug discovery
- ✓ Discovery of genetic sequences linked to diseases
- ✓ Medical imaging diagnosis
- ✓ Identifying disease and diagnosis

Government/Politics

- ✓ Police work (crime prediction)
- ✓ Election results prediction
- ✓ Traffic management

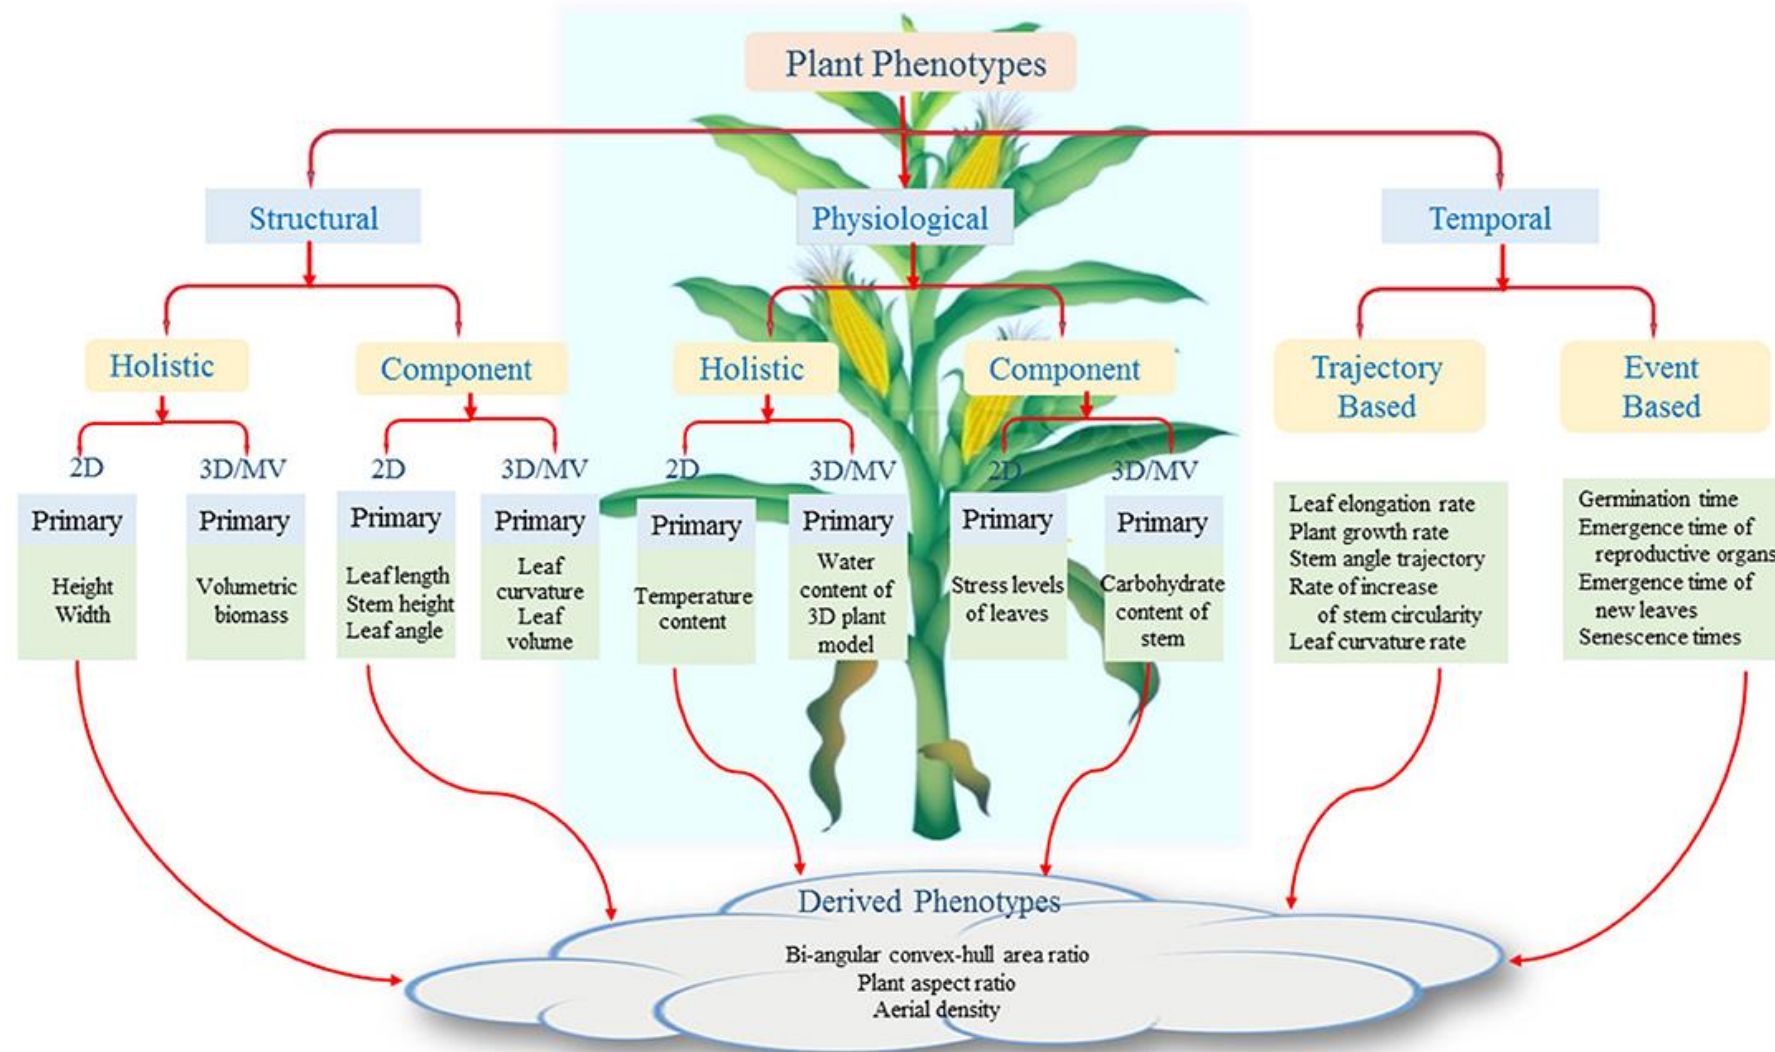
Manufacturing/BioTech

- ✓ Smart manufacturing
- ✓ Product design

And lots more!!!

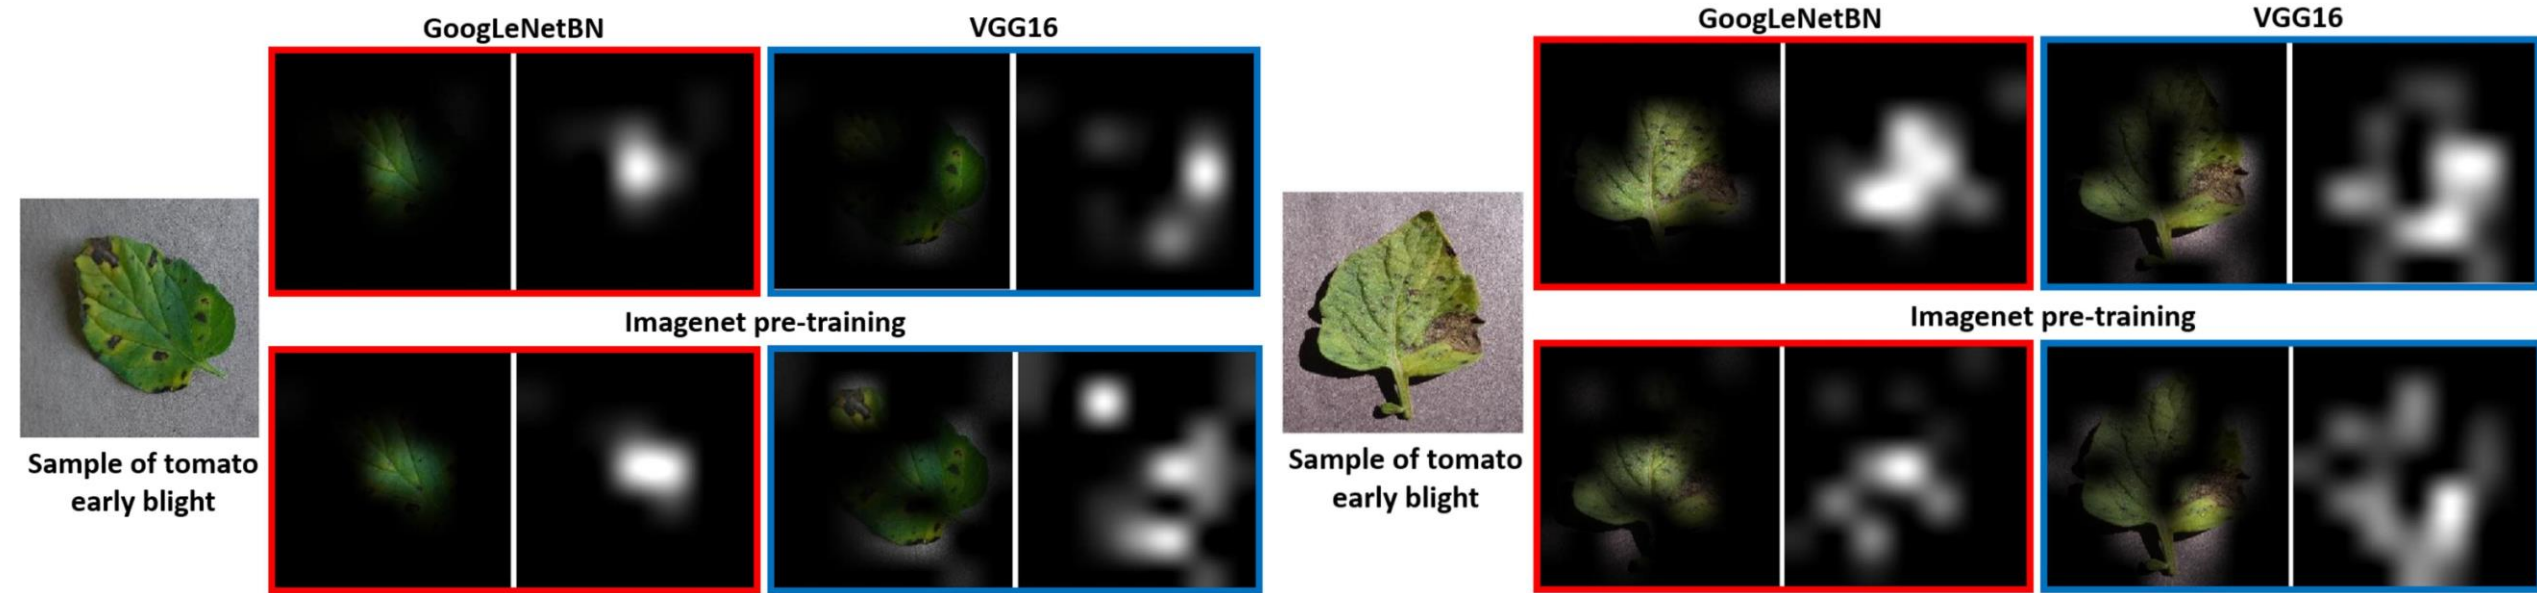
Examples of ML in plant science

High throughput phenomics



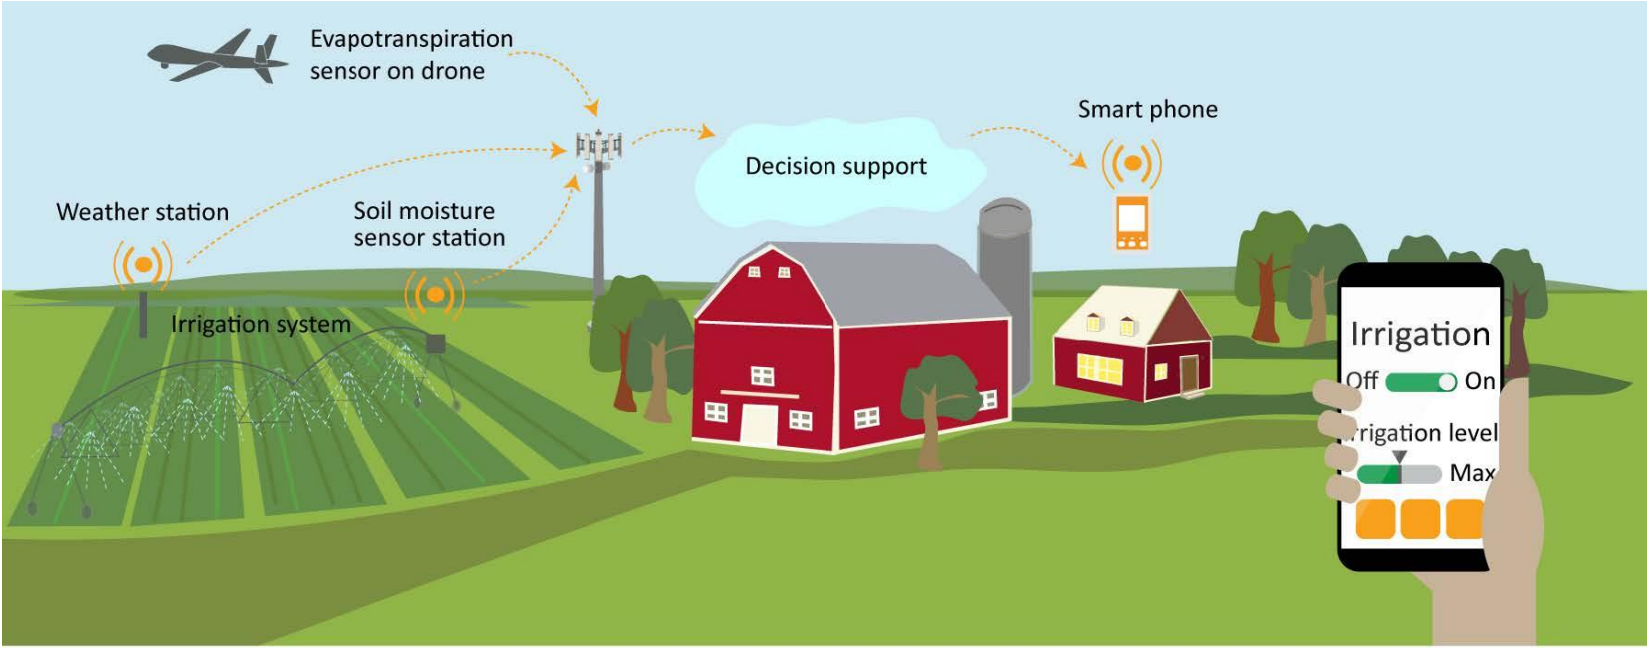
Examples of ML in plant science

Plant disease recognition and characterization



Examples of ML in plant science

Precision agriculture



Crop yield prediction

	Tester 1	Tester 2	Tester 3
Inbred 1	X	X	X
Inbred 2	X	?	?
Inbred 3	?	X	X
Inbred 4	?	?	X
Inbred 5	X	X	?

Quiz: supervised or unsupervised?

Spam detection

supervised

Anomaly detection

unsupervised

Handwriting recognition


supervised

Object recognition

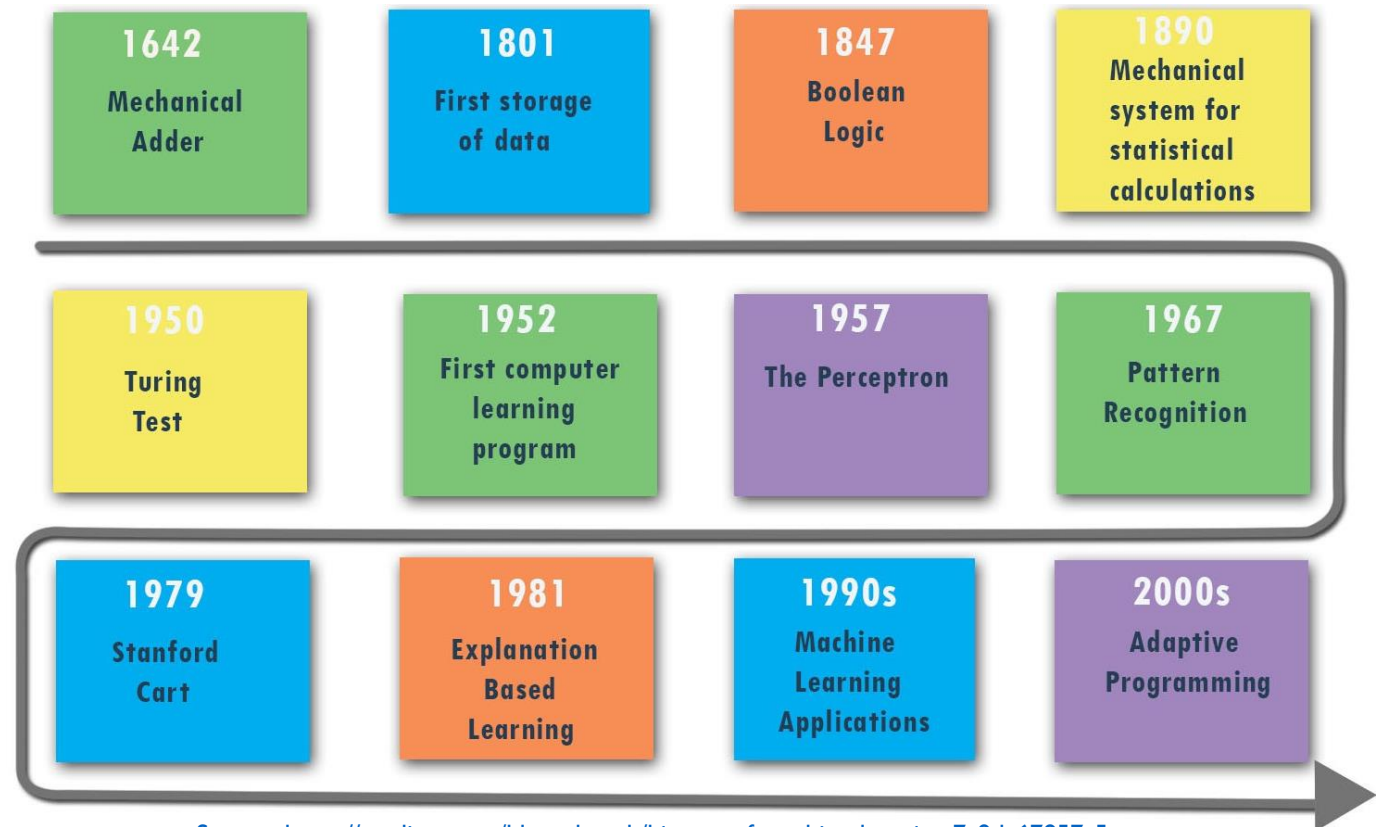
Unsupervised/supervised

History of machine learning

Cycle of advancement



Brief history




Source: <https://medium.com/bloombench/history-of-machine-learning-7c9dc67857a5>

Explore the history of machine learning with Google:


<https://cloud.withgoogle.com/build/data-analytics/explore-history-machine-learning/>

Current landscape

Total estimated investments in AI startups (2011-2017)



Current landscape



Lab 0: Laptop setup









Reproducible science

Question: Is reproducibility the same as replicability?

- **Replication:** the confirmation of results and conclusions from one study obtained independently in another
- **Reproducibility:** minimum standard for assessing the value of scientific claims, particularly when full independent replication of a study is not feasible

We want to move from this ...

Does this look familiar?

<input type="checkbox"/> Name	Date modified	Type
 Rscript_4_21_2016.R	5/1/2016 3:03 PM	R File
 Rscript_4_22_2016a.R	5/1/2016 3:03 PM	R File
 Rscript_4_22_2016b.R	5/1/2016 3:03 PM	R File
 Rscript_4_24_2016.R	5/1/2016 3:03 PM	R File
 Rscript_final.R	5/1/2016 3:03 PM	R File
 Rscript_final_final.R	5/1/2016 3:03 PM	R File
 Rscript_really_final.R	5/1/2016 3:03 PM	R File
 Rscript_really_really_final_final.R	5/1/2016 3:03 PM	R File

to more like this ...

▼ letters	
▶ cache	Store processed data
▶ config	
▶ data	Store raw data
▶ diagnostics	
▶ doc	
▶ graphs	Store figures
▶ lib	
▶ logs	
▶ munge	Scripts to preprocess data
▶ profiling	
▶ README	
▶ reports	Notebooks/markdowns: exploratory and final results
▶ src	All scripts for analyses
▶ tests	Tests to make sure code works as it should
▶ TODO	

Setting up your workstation

☐ R/R studio

- <https://www.r-project.org/>
- <https://rstudio.com/products/rstudio/#Desktop>

☐ Git / GitHub

- <https://help.github.com/en/github/getting-started-with-githubhttps://desktop.github.com/>
- <https://towardsdatascience.com/getting-started-with-git-and-github-6fcd0f2d4ac6>
- <https://happygitwithr.com/index.html>

☐ Jupyter lab/ notebooks

- <https://www.storybench.org/install-r-jupyter-notebook/>
- <https://bence.ferdinandy.com/2018/11/22/fastest-way-to-set-up-r-with-jupyter-on-windows/>

☐ Data science cookie cutters

- R packages (<http://r-pkgs.had.co.nz/>)
- R Project template (<http://projecttemplate.net/index.html>)
- Tools for writing reproducible research in R (<https://github.com/benmarwick/rrtools>)

☐ Package management (packrat: <https://rstudio.github.io/packrat/>)